

# Dhruv Bhutani

☎ 437-982-5258

✉ [dbhutani@uwaterloo.ca](mailto:dbhutani@uwaterloo.ca)

🌐 [linkedin.com/dhruvbhutani](https://www.linkedin.com/in/dhruvbhutani)

🌐 [dhruvbhutani.me](https://dhruvbhutani.me)

## Education

### The University Of Waterloo

Expected May 2025

Bachelor of Science in Computer Science

Waterloo, ON

- **Relevant Coursework:** Machine Learning, Computational Vision, Database Management & Statistics in CS

## Technical Skills

**Languages:** Python, SQL, Kotlin, JavaScript, HTML5, CSS3    **ML & AI:** PyTorch, TensorFlow, XGBoost, Hugging Face

**Databases & Cloud:** MySQL, AWS (S3, SageMaker, DynamoDB), Databricks    **Tools:** Git, Docker, Selenium, FastAPI, Flask

## Experience

### XE.com

Sept 2024 – Dec 2024

Machine Learning Engineer Intern

- **Engineered a scalable, parallel geospatial pipeline** using AWS Lambda and Pandas to process distance features across **1M+ global locations**, cutting feature engineering time by 40%.
- **Trained and optimized XGBoost models for fraud detection** on 200K+ transactions; applied feature selection and Bayesian hyperparameter tuning to boost AUC by 15%.

### Mistplay

May 2024 – Aug 2024

Machine Learning Engineer Intern

- **Built a personalized recommendation engine** combining Matrix Factorization with Hugging Face Transformer-based models, driving an estimated **10% increase in user revenue**.
- **Deployed scalable ML-driven pricing strategies** on Databricks using PySpark and SQL to automate pricing across millions of in-app transactions.
- **Created 30+ interactive dashboards and visualizations** with Plotly and Matplotlib, uncovering user behavior trends that led to **two new product features**.

### XE.com

Sept 2023 – Dec 2023

Applied ML Intern – Fraud Detection

- **Designed hybrid ML pipelines** using PyTorch, XGBoost, and AdaBoost for fraud prediction, improving detection for cross-border transactions by 8%.
- **Executed data integrity analysis** on **375K+ transaction records**, correcting systemic discrepancies to improve downstream analytics.
- **Built end-to-end ML infrastructure** using AWS SageMaker, S3, and DynamoDB, reducing inference latency and data preprocessing time by 30%.

### Teranet Inc.

Jan 2023 – Apr 2023

ML Engineer Intern – Real Estate Analytics

- **Trained ensemble models (Random Forest, XGBoost)** to predict property values with **87% accuracy**, outperforming internal benchmarks by 5%.
- **Streamlined data preprocessing workflows** using Scikit-learn and Pandas, reducing ETL times by 25% while maintaining high data quality.

### BCN Labs, Bain & Company

May 2022 – Aug 2022

AI Research Intern – Forecasting Systems

- **Developed a no-code forecasting platform** using TensorFlow and Streamlit, enabling non-technical teams to build and deploy time series models.
- **Evaluated multiple deep learning architectures** (ANNs, CNN-LSTMs, Transformers) for business forecasting, achieving up to **95% accuracy** on structured datasets.

## Projects

**MediSync: AI-Powered Clinical Research Assistant Agents** | Python, LangChain, Hugging Face Transformers, Pinecone, Streamlit

- Implemented a **chain-of-agents system** for clinicians to automate literature synthesis, note drafting, and clinical trial matching.
- Used **LLMs + Pinecone vector DB** to retrieve context-specific medical evidence, ensuring accurate and traceable AI-generated outputs.
- Reduced research and documentation time by over **50%** in simulated workflows with realistic patient case inputs.